

# DLP\_Project\_ResNet

May 12, 2025

```
[8]: !pip install tensorflow numpy pandas opencv-python matplotlib scikit-learn  
      ↪ kagglehub
```

```
Requirement already satisfied: tensorflow in /usr/local/lib/python3.11/dist-  
packages (2.18.0)  
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages  
(2.0.2)  
Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages  
(2.2.2)  
Requirement already satisfied: opencv-python in /usr/local/lib/python3.11/dist-  
packages (4.11.0.86)  
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-  
packages (3.10.0)  
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.11/dist-  
packages (1.6.1)  
Requirement already satisfied: kagglehub in /usr/local/lib/python3.11/dist-  
packages (0.3.12)  
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.11/dist-  
packages (from tensorflow) (1.4.0)  
Requirement already satisfied: astunparse>=1.6.0 in  
/usr/local/lib/python3.11/dist-packages (from tensorflow) (1.6.3)  
Requirement already satisfied: flatbuffers>=24.3.25 in  
/usr/local/lib/python3.11/dist-packages (from tensorflow) (25.2.10)  
Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in  
/usr/local/lib/python3.11/dist-packages (from tensorflow) (0.6.0)  
Requirement already satisfied: google-pasta>=0.1.1 in  
/usr/local/lib/python3.11/dist-packages (from tensorflow) (0.2.0)  
Requirement already satisfied: libclang>=13.0.0 in  
/usr/local/lib/python3.11/dist-packages (from tensorflow) (18.1.1)  
Requirement already satisfied: opt-einsum>=2.3.2 in  
/usr/local/lib/python3.11/dist-packages (from tensorflow) (3.4.0)  
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-  
packages (from tensorflow) (24.2)  
Requirement already satisfied:  
protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<6.0.0dev,>=3.20.3  
in /usr/local/lib/python3.11/dist-packages (from tensorflow) (5.29.4)  
Requirement already satisfied: requests<3,>=2.21.0 in  
/usr/local/lib/python3.11/dist-packages (from tensorflow) (2.32.3)
```

Requirement already satisfied: setuptools in /usr/local/lib/python3.11/dist-packages (from tensorflow) (75.2.0)  
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.0)  
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.1.0)  
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (4.13.2)  
Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.2)  
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.71.0)  
Requirement already satisfied: tensorboard<2.19,>=2.18 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.18.0)  
Requirement already satisfied: keras>=3.5.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.8.0)  
Requirement already satisfied: h5py>=3.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.13.0)  
Requirement already satisfied: ml-dtypes<0.5.0,>=0.4.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.4.1)  
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.37.1)  
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas) (2.9.0.post0)  
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)  
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)  
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3.2)  
Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (0.12.1)  
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (4.57.0)  
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.4.8)  
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (11.2.1)  
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (3.2.3)  
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.15.2)  
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.4.2)  
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (3.6.0)  
Requirement already satisfied: pyyaml in /usr/local/lib/python3.11/dist-packages (from kagglehub) (6.0.2)

Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from kagglehub) (4.67.1)

Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.11/dist-packages (from astunparse>=1.6.0->tensorflow) (0.45.1)

Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (13.9.4)

Requirement already satisfied: namex in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.0.9)

Requirement already satisfied: optree in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.15.0)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.4.1)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.10)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (2.4.0)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (2025.4.26)

Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (3.8)

Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (0.7.2)

Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (3.1.3)

Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.11/dist-packages (from werkzeug>=1.0.1->tensorboard<2.19,>=2.18->tensorflow) (3.0.2)

Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow) (3.0.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow) (2.19.1)

Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich->keras>=3.5.0->tensorflow) (0.1.2)

```
[9]: import os
import cv2
import numpy as np
import random
```

```

import matplotlib.pyplot as plt
from tqdm import tqdm
from collections import Counter

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.utils.class_weight import compute_class_weight

import tensorflow as tf
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.applications.resnet50 import preprocess_input
from tensorflow.keras.layers import Conv2D, Input, Dense, Dropout, _
    ↪ GlobalAveragePooling2D, Lambda, BatchNormalization, MaxPooling2D, Flatten
from tensorflow.keras.models import Model
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from tensorflow.keras import backend as K

```

```

[10]: import kagglehub
path = kagglehub.dataset_download("tejasreddy/iam-handwriting-top50")
print("Path to dataset files:", path)

```

Path to dataset files: /kaggle/input/iam-handwriting-top50

```

[11]: image_dir = '/kaggle/input/iam-handwriting-top50/data_subset/data_subset'
txt_path = '/kaggle/input/iam-handwriting-top50/forms_for_parsing.txt'

```

```

[12]: def preprocess_image(image_path, target_size=(224,224)):
    img = cv2.imread(image_path)
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img = cv2.resize(img, target_size)
    img = preprocess_input(img.astype(np.float32))
    return img

```

```

[13]: form_writer_map = {}
with open(txt_path, 'r') as f:
    for line in f:
        if line.strip():
            parts = line.strip().split()
            if len(parts) >= 2:
                form_id, writer_id = parts[0], parts[1]
                form_writer_map[form_id] = writer_id

```

```

[7]: data, labels = [], []
for fname in tqdm(os.listdir(image_dir)):
    if fname.lower().endswith(('png', '.jpg', '.jpeg')):

```

```

        form_id = os.path.splitext(fname)[0].split("-s")[0]
        if form_id in form_writer_map:
            label = form_writer_map[form_id]
            img_path = os.path.join(image_dir, fname)
            img = preprocess_image(img_path)
            data.append(img)
            labels.append(label)
X = np.array(data)
y = np.array(labels)

```

100%| | 4899/4899 [00:41<00:00, 117.36it/s]

```

[14]: from collections import Counter
label_counts = Counter(labels)
top_writers = [writer for writer, count in label_counts.most_common(10)]

filtered_data = []
filtered_labels = []
for img, label in zip(data, labels):
    if label in top_writers:
        filtered_data.append(img)
        filtered_labels.append(label)

X = np.array(filtered_data)
y = np.array(filtered_labels)

```

```

[15]: le = LabelEncoder()
y_encoded = le.fit_transform(y)
y_categorical = to_categorical(y_encoded)

```

```

[16]: X_train, X_test, y_train, y_test = train_test_split(X, y_categorical,
    ↪test_size=0.2, random_state=42)

```

```


[17]: import matplotlib.pyplot as plt

for i in range(10):
    plt.imshow(X_train[i].squeeze(), cmap='gray')
    label = le.inverse_transform([np.argmax(y_train[i])])[0]
    plt.title(f"Label: {label}")
    plt.axis('off')
    plt.show()

```


WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-87.68..151.061].

Label: 670




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-63.68..151.061].

Label: 551



WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-99.68..151.061].


Label: 552



WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-86.68..151.061].




Label: 153




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-89.68..151.061].

Label: 670



WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-70.68..151.061].

Label: 384




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-86.68..151.061].

My Mother's friend

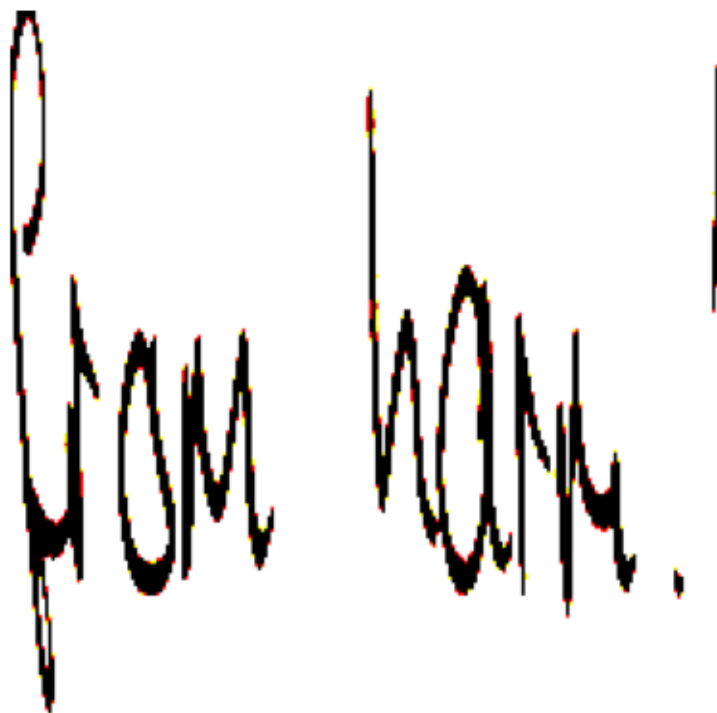
```
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with
RGB data ([0..1] for floats or [0..255] for integers). Got range
[-90.68..151.061].
```

Label: 000



WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-100.68..151.061].


Label: 552



from now!

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-76.68..151.061].

Label: 154



```
[18]: base_model = ResNet50(weights='imagenet', include_top=False,
    ↪input_shape=(224,224,3))
base_model.trainable = False # Freeze base layers initially

x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dropout(0.5)(x)
x = Dense(128, activation='relu')(x)
output = Dense(len(le.classes_), activation='softmax')(x)

resnet_model = Model(inputs=base_model.input, outputs=output)
resnet_model.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])
resnet_model.summary()

callbacks = [
    EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True),
    ReduceLROnPlateau(monitor='val_loss', patience=2, factor=0.5)
]
```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5)

94765736/94765736

1s

0us/step

Model: "functional"

Layer (type)	Output Shape	Param #	Connected to
input_layer (InputLayer)	(None, 224, 224, 3)	0	-
conv1_pad (ZeroPadding2D)	(None, 230, 230, 3)	0	input_layer[0][0]
conv1_conv (Conv2D)	(None, 112, 112, 64)	9,472	conv1_pad[0][0]
conv1_bn (BatchNormalizatio...	(None, 112, 112, 64)	256	conv1_conv[0][0]
conv1_relu (Activation)	(None, 112, 112, 64)	0	conv1_bn[0][0]
pool1_pad (ZeroPadding2D)	(None, 114, 114, 64)	0	conv1_relu[0][0]
pool1_pool (MaxPooling2D)	(None, 56, 56, 64)	0	pool1_pad[0][0]
conv2_block1_1_conv (Conv2D)	(None, 56, 56, 64)	4,160	pool1_pool[0][0]
conv2_block1_1_bn (BatchNormalizatio...	(None, 56, 56, 64)	256	conv2_block1_1_c...
conv2_block1_1_relu (Activation)	(None, 56, 56, 64)	0	conv2_block1_1_b...
conv2_block1_2_conv (Conv2D)	(None, 56, 56, 64)	36,928	conv2_block1_1_r...
conv2_block1_2_bn (BatchNormalizatio...	(None, 56, 56, 64)	256	conv2_block1_2_c...
conv2_block1_2_relu (Activation)	(None, 56, 56, 64)	0	conv2_block1_2_b...
conv2_block1_0_conv	(None, 56, 56,	16,640	pool1_pool[0][0]



(Conv2D)	256)		
conv2_block1_3_conv (Conv2D)	(None, 56, 56, 256)	16,640	conv2_block1_2_r...
conv2_block1_0_bn (BatchNormalizatio...	(None, 56, 56, 256)	1,024	conv2_block1_0_c...
conv2_block1_3_bn (BatchNormalizatio...	(None, 56, 56, 256)	1,024	conv2_block1_3_c...
conv2_block1_add (Add)	(None, 56, 56, 256)	0	conv2_block1_0_b... conv2_block1_3_b...
conv2_block1_out (Activation)	(None, 56, 56, 256)	0	conv2_block1_add...
conv2_block2_1_conv (Conv2D)	(None, 56, 56, 64)	16,448	conv2_block1_out...
conv2_block2_1_bn (BatchNormalizatio...	(None, 56, 56, 64)	256	conv2_block2_1_c...
conv2_block2_1_relu (Activation)	(None, 56, 56, 64)	0	conv2_block2_1_b...
conv2_block2_2_conv (Conv2D)	(None, 56, 56, 64)	36,928	conv2_block2_1_r...
conv2_block2_2_bn (BatchNormalizatio...	(None, 56, 56, 64)	256	conv2_block2_2_c...
conv2_block2_2_relu (Activation)	(None, 56, 56, 64)	0	conv2_block2_2_b...
conv2_block2_3_conv (Conv2D)	(None, 56, 56, 256)	16,640	conv2_block2_2_r...
conv2_block2_3_bn (BatchNormalizatio...	(None, 56, 56, 256)	1,024	conv2_block2_3_c...
conv2_block2_add (Add)	(None, 56, 56, 256)	0	conv2_block1_out... conv2_block2_3_b...
conv2_block2_out (Activation)	(None, 56, 56, 256)	0	conv2_block2_add...
conv2_block3_1_conv	(None, 56, 56,	16,448	conv2_block2_out...

(Conv2D)	64)		
conv2_block3_1_bn (BatchNormalizatio...	(None, 56, 56, 64)	256	conv2_block3_1_c...
conv2_block3_1_relu (Activation)	(None, 56, 56, 64)	0	conv2_block3_1_b...
conv2_block3_2_conv (Conv2D)	(None, 56, 56, 64)	36,928	conv2_block3_1_r...
conv2_block3_2_bn (BatchNormalizatio...	(None, 56, 56, 64)	256	conv2_block3_2_c...
conv2_block3_2_relu (Activation)	(None, 56, 56, 64)	0	conv2_block3_2_b...
conv2_block3_3_conv (Conv2D)	(None, 56, 56, 256)	16,640	conv2_block3_2_r...
conv2_block3_3_bn (BatchNormalizatio...	(None, 56, 56, 256)	1,024	conv2_block3_3_c...
conv2_block3_add (Add)	(None, 56, 56, 256)	0	conv2_block2_out... conv2_block3_3_b...
conv2_block3_out (Activation)	(None, 56, 56, 256)	0	conv2_block3_add...
conv3_block1_1_conv (Conv2D)	(None, 28, 28, 128)	32,896	conv2_block3_out...
conv3_block1_1_bn (BatchNormalizatio...	(None, 28, 28, 128)	512	conv3_block1_1_c...
conv3_block1_1_relu (Activation)	(None, 28, 28, 128)	0	conv3_block1_1_b...
conv3_block1_2_conv (Conv2D)	(None, 28, 28, 128)	147,584	conv3_block1_1_r...
conv3_block1_2_bn (BatchNormalizatio...	(None, 28, 28, 128)	512	conv3_block1_2_c...
conv3_block1_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block1_2_b...
conv3_block1_0_conv	(None, 28, 28,	131,584	conv2_block3_out...

(Conv2D)	512)		
conv3_block1_3_conv (Conv2D)	(None, 28, 28, 512)	66,048	conv3_block1_2_r...
conv3_block1_0_bn (BatchNormalizatio...	(None, 28, 28, 512)	2,048	conv3_block1_0_c...
conv3_block1_3_bn (BatchNormalizatio...	(None, 28, 28, 512)	2,048	conv3_block1_3_c...
conv3_block1_add (Add)	(None, 28, 28, 512)	0	conv3_block1_0_b... conv3_block1_3_b...
conv3_block1_out (Activation)	(None, 28, 28, 512)	0	conv3_block1_add...
conv3_block2_1_conv (Conv2D)	(None, 28, 28, 128)	65,664	conv3_block1_out...
conv3_block2_1_bn (BatchNormalizatio...	(None, 28, 28, 128)	512	conv3_block2_1_c...
conv3_block2_1_relu (Activation)	(None, 28, 28, 128)	0	conv3_block2_1_b...
conv3_block2_2_conv (Conv2D)	(None, 28, 28, 128)	147,584	conv3_block2_1_r...
conv3_block2_2_bn (BatchNormalizatio...	(None, 28, 28, 128)	512	conv3_block2_2_c...
conv3_block2_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block2_2_b...
conv3_block2_3_conv (Conv2D)	(None, 28, 28, 512)	66,048	conv3_block2_2_r...
conv3_block2_3_bn (BatchNormalizatio...	(None, 28, 28, 512)	2,048	conv3_block2_3_c...
conv3_block2_add (Add)	(None, 28, 28, 512)	0	conv3_block1_out... conv3_block2_3_b...
conv3_block2_out (Activation)	(None, 28, 28, 512)	0	conv3_block2_add...
conv3_block3_1_conv	(None, 28, 28,	65,664	conv3_block2_out...

(Conv2D)	128)		
conv3_block3_1_bn (BatchNormalizatio...	(None, 28, 28, 128)	512	conv3_block3_1_c...
conv3_block3_1_relu (Activation)	(None, 28, 28, 128)	0	conv3_block3_1_b...
conv3_block3_2_conv (Conv2D)	(None, 28, 28, 128)	147,584	conv3_block3_1_r...
conv3_block3_2_bn (BatchNormalizatio...	(None, 28, 28, 128)	512	conv3_block3_2_c...
conv3_block3_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block3_2_b...
conv3_block3_3_conv (Conv2D)	(None, 28, 28, 512)	66,048	conv3_block3_2_r...
conv3_block3_3_bn (BatchNormalizatio...	(None, 28, 28, 512)	2,048	conv3_block3_3_c...
conv3_block3_add (Add)	(None, 28, 28, 512)	0	conv3_block2_out... conv3_block3_3_b...
conv3_block3_out (Activation)	(None, 28, 28, 512)	0	conv3_block3_add...
conv3_block4_1_conv (Conv2D)	(None, 28, 28, 128)	65,664	conv3_block3_out...
conv3_block4_1_bn (BatchNormalizatio...	(None, 28, 28, 128)	512	conv3_block4_1_c...
conv3_block4_1_relu (Activation)	(None, 28, 28, 128)	0	conv3_block4_1_b...
conv3_block4_2_conv (Conv2D)	(None, 28, 28, 128)	147,584	conv3_block4_1_r...
conv3_block4_2_bn (BatchNormalizatio...	(None, 28, 28, 128)	512	conv3_block4_2_c...
conv3_block4_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block4_2_b...
conv3_block4_3_conv	(None, 28, 28,	66,048	conv3_block4_2_r...

(Conv2D)	512)		
conv3_block4_3_bn (BatchNormalizatio...	(None, 28, 28, 512)	2,048	conv3_block4_3_c...
conv3_block4_add (Add)	(None, 28, 28, 512)	0	conv3_block3_out... conv3_block4_3_b...
conv3_block4_out (Activation)	(None, 28, 28, 512)	0	conv3_block4_add...
conv4_block1_1_conv (Conv2D)	(None, 14, 14, 256)	131,328	conv3_block4_out...
conv4_block1_1_bn (BatchNormalizatio...	(None, 14, 14, 256)	1,024	conv4_block1_1_c...
conv4_block1_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block1_1_b...
conv4_block1_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block1_1_r...
conv4_block1_2_bn (BatchNormalizatio...	(None, 14, 14, 256)	1,024	conv4_block1_2_c...
conv4_block1_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block1_2_b...
conv4_block1_0_conv (Conv2D)	(None, 14, 14, 1024)	525,312	conv3_block4_out...
conv4_block1_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block1_2_r...
conv4_block1_0_bn (BatchNormalizatio...	(None, 14, 14, 1024)	4,096	conv4_block1_0_c...
conv4_block1_3_bn (BatchNormalizatio...	(None, 14, 14, 1024)	4,096	conv4_block1_3_c...
conv4_block1_add (Add)	(None, 14, 14, 1024)	0	conv4_block1_0_b... conv4_block1_3_b...
conv4_block1_out (Activation)	(None, 14, 14, 1024)	0	conv4_block1_add...
conv4_block2_1_conv	(None, 14, 14,	262,400	conv4_block1_out...

(Conv2D)	256)		
conv4_block2_1_bn (BatchNormalizatio...	(None, 14, 14, 256)	1,024	conv4_block2_1_c...
conv4_block2_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block2_1_b...
conv4_block2_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block2_1_r...
conv4_block2_2_bn (BatchNormalizatio...	(None, 14, 14, 256)	1,024	conv4_block2_2_c...
conv4_block2_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block2_2_b...
conv4_block2_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block2_2_r...
conv4_block2_3_bn (BatchNormalizatio...	(None, 14, 14, 1024)	4,096	conv4_block2_3_c...
conv4_block2_add (Add)	(None, 14, 14, 1024)	0	conv4_block1_out... conv4_block2_3_b...
conv4_block2_out (Activation)	(None, 14, 14, 1024)	0	conv4_block2_add...
conv4_block3_1_conv (Conv2D)	(None, 14, 14, 256)	262,400	conv4_block2_out...
conv4_block3_1_bn (BatchNormalizatio...	(None, 14, 14, 256)	1,024	conv4_block3_1_c...
conv4_block3_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block3_1_b...
conv4_block3_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block3_1_r...
conv4_block3_2_bn (BatchNormalizatio...	(None, 14, 14, 256)	1,024	conv4_block3_2_c...
conv4_block3_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block3_2_b...
conv4_block3_3_conv	(None, 14, 14,	263,168	conv4_block3_2_r...

(Conv2D)	1024)		
conv4_block3_3_bn (BatchNormalizatio...	(None, 14, 14, 1024)	4,096	conv4_block3_3_c...
conv4_block3_add (Add)	(None, 14, 14, 1024)	0	conv4_block2_out... conv4_block3_3_b...
conv4_block3_out (Activation)	(None, 14, 14, 1024)	0	conv4_block3_add...
conv4_block4_1_conv (Conv2D)	(None, 14, 14, 256)	262,400	conv4_block3_out...
conv4_block4_1_bn (BatchNormalizatio...	(None, 14, 14, 256)	1,024	conv4_block4_1_c...
conv4_block4_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block4_1_b...
conv4_block4_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block4_1_r...
conv4_block4_2_bn (BatchNormalizatio...	(None, 14, 14, 256)	1,024	conv4_block4_2_c...
conv4_block4_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block4_2_b...
conv4_block4_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block4_2_r...
conv4_block4_3_bn (BatchNormalizatio...	(None, 14, 14, 1024)	4,096	conv4_block4_3_c...
conv4_block4_add (Add)	(None, 14, 14, 1024)	0	conv4_block3_out... conv4_block4_3_b...
conv4_block4_out (Activation)	(None, 14, 14, 1024)	0	conv4_block4_add...
conv4_block5_1_conv (Conv2D)	(None, 14, 14, 256)	262,400	conv4_block4_out...
conv4_block5_1_bn (BatchNormalizatio...	(None, 14, 14, 256)	1,024	conv4_block5_1_c...
conv4_block5_1_relu	(None, 14, 14,	0	conv4_block5_1_b...

(Activation)	256)		
conv4_block5_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block5_1_r...
conv4_block5_2_bn (BatchNormalizatio...	(None, 14, 14, 256)	1,024	conv4_block5_2_c...
conv4_block5_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block5_2_b...
conv4_block5_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block5_2_r...
conv4_block5_3_bn (BatchNormalizatio...	(None, 14, 14, 1024)	4,096	conv4_block5_3_c...
conv4_block5_add (Add)	(None, 14, 14, 1024)	0	conv4_block4_out... conv4_block5_3_b...
conv4_block5_out (Activation)	(None, 14, 14, 1024)	0	conv4_block5_add...
conv4_block6_1_conv (Conv2D)	(None, 14, 14, 256)	262,400	conv4_block5_out...
conv4_block6_1_bn (BatchNormalizatio...	(None, 14, 14, 256)	1,024	conv4_block6_1_c...
conv4_block6_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block6_1_b...
conv4_block6_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block6_1_r...
conv4_block6_2_bn (BatchNormalizatio...	(None, 14, 14, 256)	1,024	conv4_block6_2_c...
conv4_block6_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block6_2_b...
conv4_block6_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block6_2_r...
conv4_block6_3_bn (BatchNormalizatio...	(None, 14, 14, 1024)	4,096	conv4_block6_3_c...
conv4_block6_add	(None, 14, 14,	0	conv4_block5_out...



(Add)	1024)		conv4_block6_3_b...
conv4_block6_out (Activation)	(None, 14, 14, 1024)	0	conv4_block6_add...
conv5_block1_1_conv (Conv2D)	(None, 7, 7, 512)	524,800	conv4_block6_out...
conv5_block1_1_bn (BatchNormalizatio...	(None, 7, 7, 512)	2,048	conv5_block1_1_c...
conv5_block1_1_relu (Activation)	(None, 7, 7, 512)	0	conv5_block1_1_b...
conv5_block1_2_conv (Conv2D)	(None, 7, 7, 512)	2,359,808	conv5_block1_1_r...
conv5_block1_2_bn (BatchNormalizatio...	(None, 7, 7, 512)	2,048	conv5_block1_2_c...
conv5_block1_2_relu (Activation)	(None, 7, 7, 512)	0	conv5_block1_2_b...
conv5_block1_0_conv (Conv2D)	(None, 7, 7, 2048)	2,099,200	conv4_block6_out...
conv5_block1_3_conv (Conv2D)	(None, 7, 7, 2048)	1,050,624	conv5_block1_2_r...
conv5_block1_0_bn (BatchNormalizatio...	(None, 7, 7, 2048)	8,192	conv5_block1_0_c...
conv5_block1_3_bn (BatchNormalizatio...	(None, 7, 7, 2048)	8,192	conv5_block1_3_c...
conv5_block1_add (Add)	(None, 7, 7, 2048)	0	conv5_block1_0_b... conv5_block1_3_b...
conv5_block1_out (Activation)	(None, 7, 7, 2048)	0	conv5_block1_add...
conv5_block2_1_conv (Conv2D)	(None, 7, 7, 512)	1,049,088	conv5_block1_out...
conv5_block2_1_bn (BatchNormalizatio...	(None, 7, 7, 512)	2,048	conv5_block2_1_c...
conv5_block2_1_relu	(None, 7, 7, 512)	0	conv5_block2_1_b...

(Activation)			
conv5_block2_2_conv (Conv2D)	(None, 7, 7, 512)	2,359,808	conv5_block2_1_r...
conv5_block2_2_bn (BatchNormalizatio...	(None, 7, 7, 512)	2,048	conv5_block2_2_c...
conv5_block2_2_relu (Activation)	(None, 7, 7, 512)	0	conv5_block2_2_b...
conv5_block2_3_conv (Conv2D)	(None, 7, 7, 2048)	1,050,624	conv5_block2_2_r...
conv5_block2_3_bn (BatchNormalizatio...	(None, 7, 7, 2048)	8,192	conv5_block2_3_c...
conv5_block2_add (Add)	(None, 7, 7, 2048)	0	conv5_block1_out... conv5_block2_3_b...
conv5_block2_out (Activation)	(None, 7, 7, 2048)	0	conv5_block2_add...
conv5_block3_1_conv (Conv2D)	(None, 7, 7, 512)	1,049,088	conv5_block2_out...
conv5_block3_1_bn (BatchNormalizatio...	(None, 7, 7, 512)	2,048	conv5_block3_1_c...
conv5_block3_1_relu (Activation)	(None, 7, 7, 512)	0	conv5_block3_1_b...
conv5_block3_2_conv (Conv2D)	(None, 7, 7, 512)	2,359,808	conv5_block3_1_r...
conv5_block3_2_bn (BatchNormalizatio...	(None, 7, 7, 512)	2,048	conv5_block3_2_c...
conv5_block3_2_relu (Activation)	(None, 7, 7, 512)	0	conv5_block3_2_b...
conv5_block3_3_conv (Conv2D)	(None, 7, 7, 2048)	1,050,624	conv5_block3_2_r...
conv5_block3_3_bn (BatchNormalizatio...	(None, 7, 7, 2048)	8,192	conv5_block3_3_c...
conv5_block3_add	(None, 7, 7,	0	conv5_block2_out...

(Add)	2048)		conv5_block3_3_b...
conv5_block3_out (Activation)	(None, 7, 7, 2048)	0	conv5_block3_add...
global_average_poo... (GlobalAveragePool...	(None, 2048)	0	conv5_block3_out...
dropout (Dropout)	(None, 2048)	0	global_average_p...
dense (Dense)	(None, 128)	262,272	dropout[0][0]
dense_1 (Dense)	(None, 10)	1,290	dense[0][0]

Total params: 23,851,274 (90.99 MB)

Trainable params: 263,562 (1.01 MB)

Non-trainable params: 23,587,712 (89.98 MB)

```
[19]: history = resnet_model.fit(
    X_train, y_train,
    validation_data=(X_test, y_test),
    batch_size=32,
    epochs=15,
    callbacks=callbacks
)
```

Epoch 1/15

45/45 31s 383ms/step -

accuracy: 0.4446 - loss: 1.8405 - val\_accuracy: 0.7827 - val\_loss: 0.7274 -  
learning\_rate: 0.0010

Epoch 2/15

45/45 5s 101ms/step -

accuracy: 0.7313 - loss: 0.7936 - val\_accuracy: 0.8914 - val\_loss: 0.3952 -  
learning\_rate: 0.0010

Epoch 3/15

45/45 5s 101ms/step -

accuracy: 0.8180 - loss: 0.5350 - val\_accuracy: 0.9164 - val\_loss: 0.3032 -  
learning\_rate: 0.0010

Epoch 4/15

45/45 5s 103ms/step -

accuracy: 0.8611 - loss: 0.3947 - val\_accuracy: 0.9248 - val\_loss: 0.2602 -  
learning\_rate: 0.0010

```

Epoch 5/15
45/45          5s 102ms/step -
accuracy: 0.8976 - loss: 0.3119 - val_accuracy: 0.9192 - val_loss: 0.2476 -
learning_rate: 0.0010
Epoch 6/15
45/45          6s 113ms/step -
accuracy: 0.9065 - loss: 0.2968 - val_accuracy: 0.9359 - val_loss: 0.2276 -
learning_rate: 0.0010
Epoch 7/15
45/45          5s 111ms/step -
accuracy: 0.9064 - loss: 0.2775 - val_accuracy: 0.9499 - val_loss: 0.1921 -
learning_rate: 0.0010
Epoch 8/15
45/45          5s 106ms/step -
accuracy: 0.9322 - loss: 0.1954 - val_accuracy: 0.9499 - val_loss: 0.1894 -
learning_rate: 0.0010
Epoch 9/15
45/45          5s 109ms/step -
accuracy: 0.9431 - loss: 0.1825 - val_accuracy: 0.9415 - val_loss: 0.2185 -
learning_rate: 0.0010
Epoch 10/15
45/45          5s 101ms/step -
accuracy: 0.9364 - loss: 0.1771 - val_accuracy: 0.9331 - val_loss: 0.2220 -
learning_rate: 0.0010
Epoch 11/15
45/45          5s 105ms/step -
accuracy: 0.9494 - loss: 0.1420 - val_accuracy: 0.9582 - val_loss: 0.1555 -
learning_rate: 5.0000e-04
Epoch 12/15
45/45          5s 110ms/step -
accuracy: 0.9557 - loss: 0.1369 - val_accuracy: 0.9526 - val_loss: 0.1759 -
learning_rate: 5.0000e-04
Epoch 13/15
45/45          5s 113ms/step -
accuracy: 0.9537 - loss: 0.1389 - val_accuracy: 0.9526 - val_loss: 0.1870 -
learning_rate: 5.0000e-04
Epoch 14/15
45/45          5s 111ms/step -
accuracy: 0.9613 - loss: 0.1226 - val_accuracy: 0.9554 - val_loss: 0.1752 -
learning_rate: 2.5000e-04
Epoch 15/15
45/45          5s 116ms/step -
accuracy: 0.9724 - loss: 0.0929 - val_accuracy: 0.9666 - val_loss: 0.1462 -
learning_rate: 2.5000e-04


```

```


[20]: import matplotlib.pyplot as plt
      plt.plot(history.history['accuracy'], label='Train')

```

```
plt.plot(history.history['val_accuracy'], label='Test')
plt.title('Accuracy over epochs')
plt.legend()
plt.show()
```



```
[21]: import matplotlib.pyplot as plt
plt.plot(history.history['loss'], label='Train')
plt.plot(history.history['val_loss'], label='Test')
plt.title('Loss over epochs')
plt.legend()
plt.show()
```



```
[22]: resnet_model.save("/kaggle/resnet_model.keras")
```


```
[23]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import numpy as np

y_pred = resnet_model.predict(X_test)
y_pred_labels = np.argmax(y_pred, axis=1)
y_true_labels = np.argmax(y_test, axis=1)

cm = confusion_matrix(y_true_labels, y_pred_labels)
disp = ConfusionMatrixDisplay(cm, display_labels=le.classes_)
disp.plot(cmap='viridis', xticks_rotation=90)
```

12/12                      8s 393ms/step

```
[23]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7abc62514710>
```




```
[24]: preds = resnet_model.predict(X_test)
for i in range(5):
    plt.imshow(X_test[i].squeeze(), cmap='gray')
    true_label = le.inverse_transform([np.argmax(y_test[i])])[0]
    pred_label = le.inverse_transform([np.argmax(preds[i])])[0]
    plt.title(f"True: {true_label} | Pred: {pred_label}")
    plt.axis('off')
    plt.show()
```

12/12 1s 80ms/step

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-72.68..151.061].


True: 671 | Pred: 671



WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-89.68..151.061].




True: 670 | Pred: 670



WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-74.68..151.061].

True: 000 | Pred: 000



WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-61.68..151.061].

True: 671 | Pred: 671

A handwritten sample of the name "Emily Catherine" in black ink. Below the name is a large, stylized "0" or "O". The entire text is overlaid with a red and yellow bounding box, indicating a segmentation or classification result. The text is written on a white background.

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-100.68..151.061].

True: 552 | Pred: 552

